

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
```

```
In [2]: df=pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv')
```

```
In [3]: df
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500	NaN	Q

891 rows × 12 columns

```
In [4]: df.columns
```

```
Out[4]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
       'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
       dtype='object')
```

```
In [5]: df.head(5)
```

Out[5]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th... Jacques Heath (Lily May Peel)	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

In [6]: df.sample(5)

Out[6]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
420	421	0	3	Gheorgheff, Mr. Stanio	male	NaN	0	0	349254	7.8958	NaN	C
749	750	0	3	Connaghton, Mr. Michael	male	31.0	0	0	335097	7.7500	NaN	Q
652	653	0	3	Kalvik, Mr. Johannes Halvorsen	male	21.0	0	0	8475	8.4333	NaN	S
156	157	1	3	Gilnagh, Miss. Katherine "Katie"	female	16.0	0	0	35851	7.7333	NaN	Q
575	576	0	3	Patchett, Mr. George	male	19.0	0	0	358585	14.5000	NaN	S

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   PassengerId  891 non-null    int64  
 1   Survived     891 non-null    int64  
 2   Pclass       891 non-null    int64  
 3   Name         891 non-null    object  
 4   Sex          891 non-null    object  
 5   Age          714 non-null    float64 
 6   SibSp        891 non-null    int64  
 7   Parch        891 non-null    int64  
 8   Ticket       891 non-null    object  
 9   Fare          891 non-null    float64 
 10  Cabin        204 non-null    object  
 11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [8]: df.describe()

Out[8]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [9]: `df.describe(include="object")`

Out[9]:

	Name	Sex	Ticket	Cabin	Embarked
count	891	891	891	204	889
unique	891	2	681	147	3
top	Dooley, Mr. Patrick	male	347082	G6	S
freq	1	577	7	4	644

In [10]: `#Drop Unnecessary Columns`
`df1=df.drop(["PassengerId", "Cabin", "Name", "Ticket"], axis=1)`
`df1`

Out[10]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S
...
886	0	2	male	27.0	0	0	13.0000	S
887	1	1	female	19.0	0	0	30.0000	S
888	0	3	female	Nan	1	2	23.4500	S
889	1	1	male	26.0	0	0	30.0000	C
890	0	3	male	32.0	0	0	7.7500	Q

891 rows × 8 columns

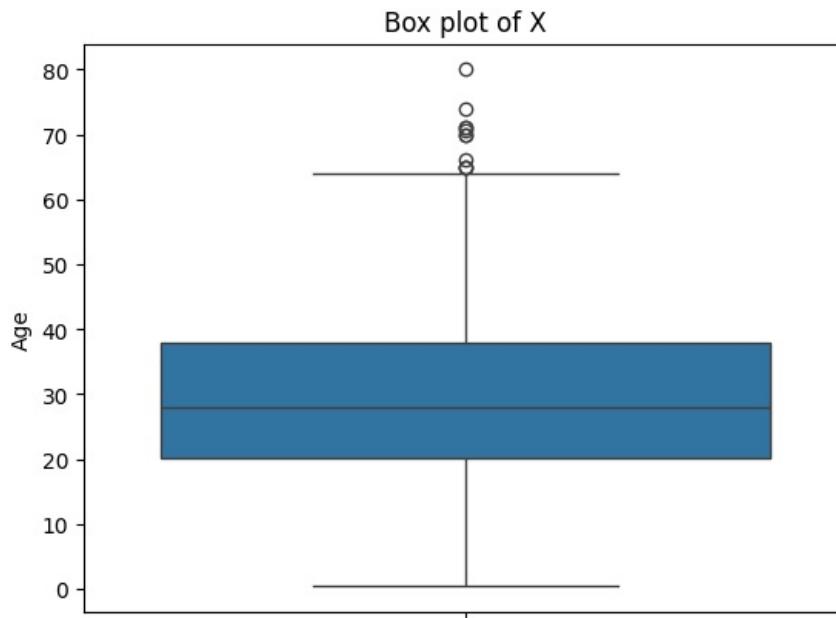
In [11]: `df1.isna().sum()`

Out[11]:

```
Survived      0
Pclass        0
Sex           0
Age          177
SibSp         0
Parch         0
Fare          0
Embarked      2
dtype: int64
```

In [12]: `#Check Age outliers`

```
sns.boxplot(df1["Age"])
plt.title("Box plot of X")
plt.show()
```



```
In [13]: df1["Embarked"].value_counts()
```

```
Out[13]: Embarked
S    644
C    168
Q     77
Name: count, dtype: int64
```

```
In [14]: df1["Age"]
```

```
Out[14]: 0      22.0
1      38.0
2      26.0
3      35.0
4      35.0
...
886    27.0
887    19.0
888    NaN
889    26.0
890    32.0
Name: Age, Length: 891, dtype: float64
```

```
In [15]: #Handling Null Values with median and mode
```

```
df1["Age"] = df1["Age"].fillna(df["Age"].median())
df1["Embarked"] = df1["Embarked"].fillna(df["Embarked"].mode()[0])
df1
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S
...
886	0	2	male	27.0	0	0	13.0000	S
887	1	1	female	19.0	0	0	30.0000	S
888	0	3	female	28.0	1	2	23.4500	S
889	1	1	male	26.0	0	0	30.0000	C
890	0	3	male	32.0	0	0	7.7500	Q

891 rows × 8 columns

```
In [16]: df1.isna().sum()
```

```
Out[16]: Survived      0  
Pclass        0  
Sex           0  
Age           0  
SibSp         0  
Parch         0  
Fare          0  
Embarked      0  
dtype: int64
```

```
In [17]: #Feature Binning
```

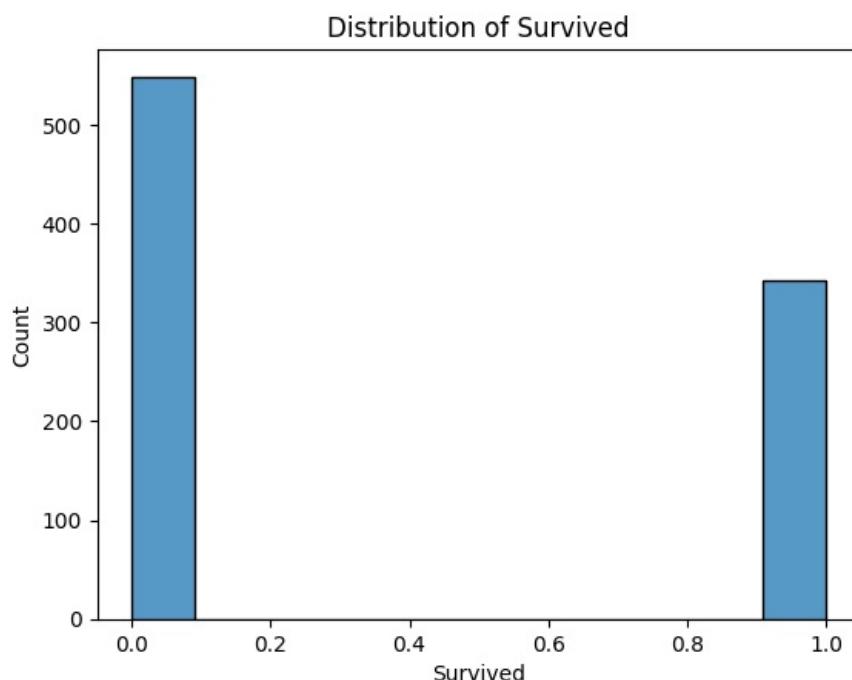
```
df1["Age_group"] = pd.cut(df1["Age"], bins=[0, 18, 35, 65, 85], labels=["Child", "Adult", "Senior", "Most Senior"])  
df1
```

```
Out[17]:   Survived  Pclass  Sex  Age  SibSp  Parch  Fare  Embarked  Age_group  
0            0       3  male  22.0     1      0  7.2500      S    Adult  
1            1       1 female  38.0     1      0  71.2833      C   Senior  
2            1       3 female  26.0     0      0  7.9250      S    Adult  
3            1       1 female  35.0     1      0  53.1000      S    Adult  
4            0       3  male  35.0     0      0  8.0500      S    Adult  
...          ...     ...  ...  ...     ...    ...  ...     ...  ...  
886           0       2  male  27.0     0      0  13.0000      S    Adult  
887           1       1 female  19.0     0      0  30.0000      S    Adult  
888           0       3 female  28.0     1      2  23.4500      S    Adult  
889           1       1  male  26.0     0      0  30.0000      C    Adult  
890           0       3  male  32.0     0      0  7.7500      Q    Adult
```

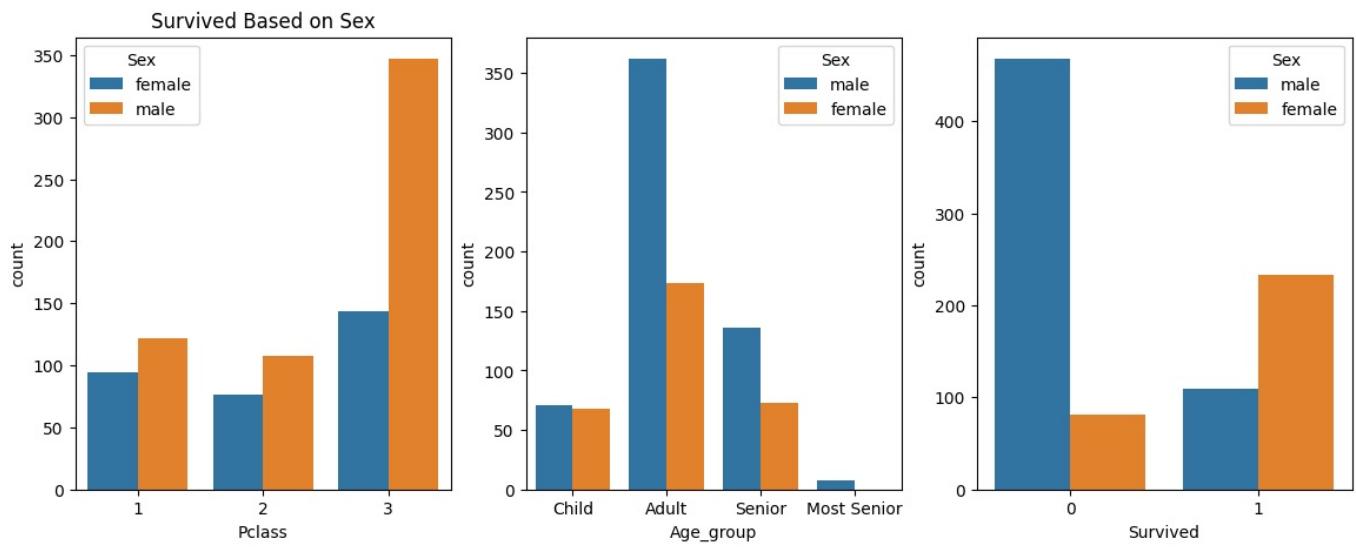
891 rows × 9 columns

```
In [18]: # Visualization of this Data
```

```
In [19]: sns.histplot(x="Survived", data=df1)  
plt.title("Distribution of Survived")  
plt.show()
```



```
In [20]: fig, ax=plt.subplots(1,3, figsize=(14,5))  
sns.countplot(x="Pclass", data=df, hue="Sex", ax=ax[0])  
sns.countplot(x="Age_group", data=df1, hue="Sex", ax=ax[1])  
sns.countplot(x="Survived", data=df1, hue="Sex", ax=ax[2])  
ax[0].set_title("PClass Based on Sex")  
ax[1].set_title("Age Group Based on Sex")  
ax[2].set_title("Survived Based on Sex")  
plt.show()
```

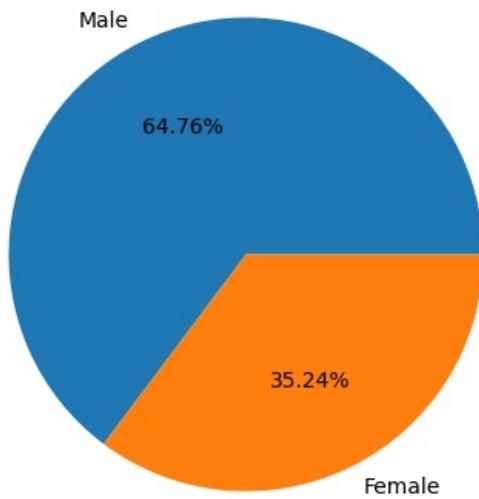


```
In [21]: sex_count=df["Sex"].value_counts()
sex_count
```

```
Out[21]: Sex
male    577
female   314
Name: count, dtype: int64
```

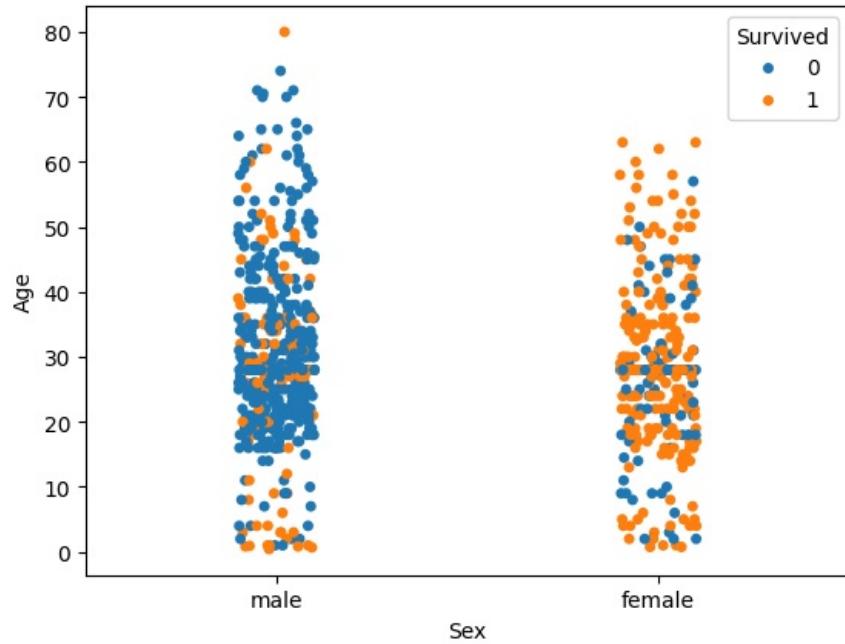
```
In [22]: plt.figure(figsize=(10,5))
plt.pie(sex_count,labels=["Male","Female"],autopct='%1.2f%%')
plt.title("Passenger Based on Gender")
plt.show()
```

Passenger Based on Gender



```
In [23]: sns.stripplot(x="Sex",y="Age",data=df1,hue="Survived")
```

```
Out[23]: <Axes: xlabel='Sex', ylabel='Age'>
```

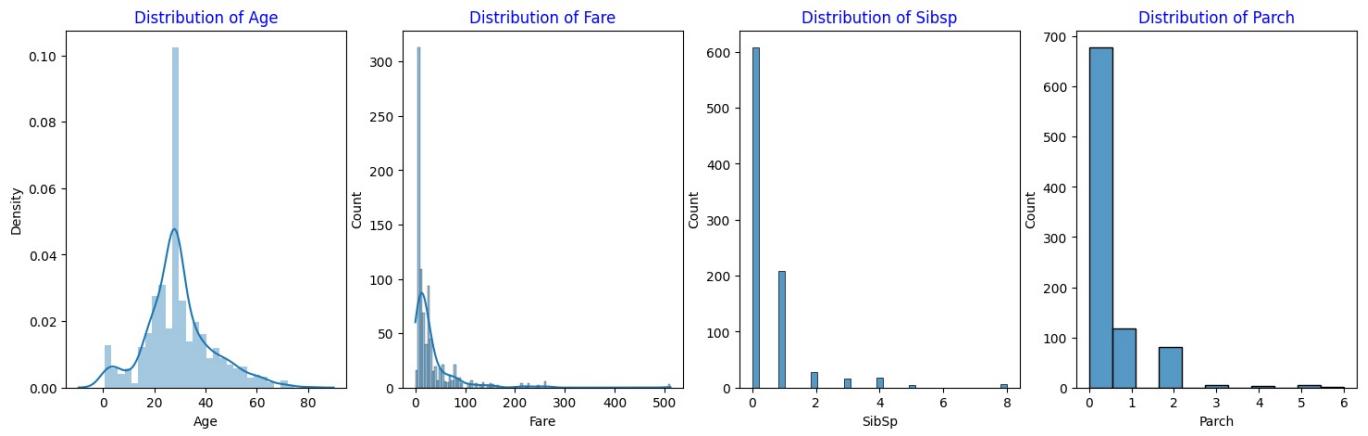


```
In [24]: fig,ax=plt.subplots(1,4,figsize=(18,5))
sns.distplot(df1["Age"],ax=ax[0])
sns.histplot(df1["Fare"],ax=ax[1],kde=True)
sns.histplot(df1["SibSp"],ax=ax[2])
sns.histplot(df1["Parch"],ax=ax[3])

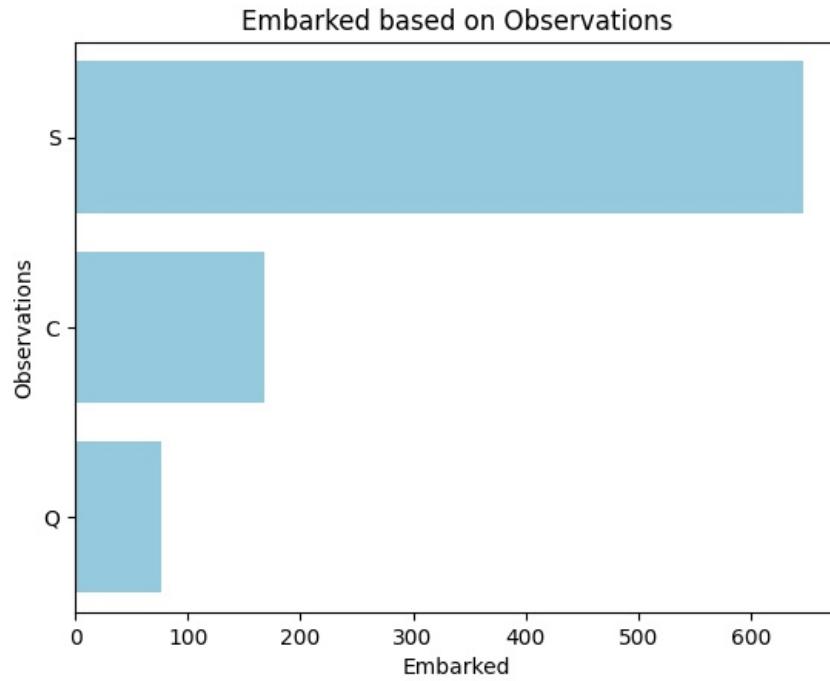
ax[0].set_title("Distribution of Age",color="Blue")
ax[1].set_title("Distribution of Fare",color="Blue")
ax[2].set_title("Distribution of Sibsp",color="Blue")
ax[3].set_title("Distribution of Parch",color="Blue")
plt.show()
```

/tmp/ipykernel_82030/2831472293.py:2: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

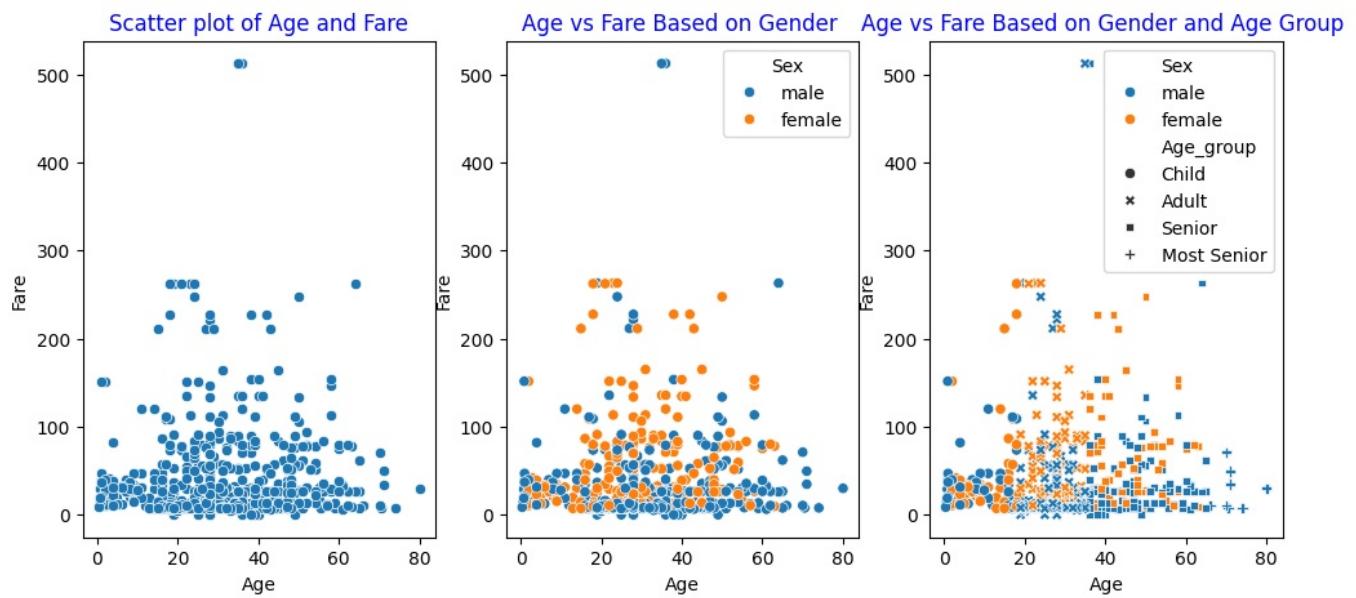
```
sns.distplot(df1["Age"],ax=ax[0])
```



```
In [25]: sns.countplot(df1["Embarked"], color="skyblue")
plt.title("Embarked based on Observations")
plt.xlabel("Embarked")
plt.ylabel("Observations")
plt.show()
```

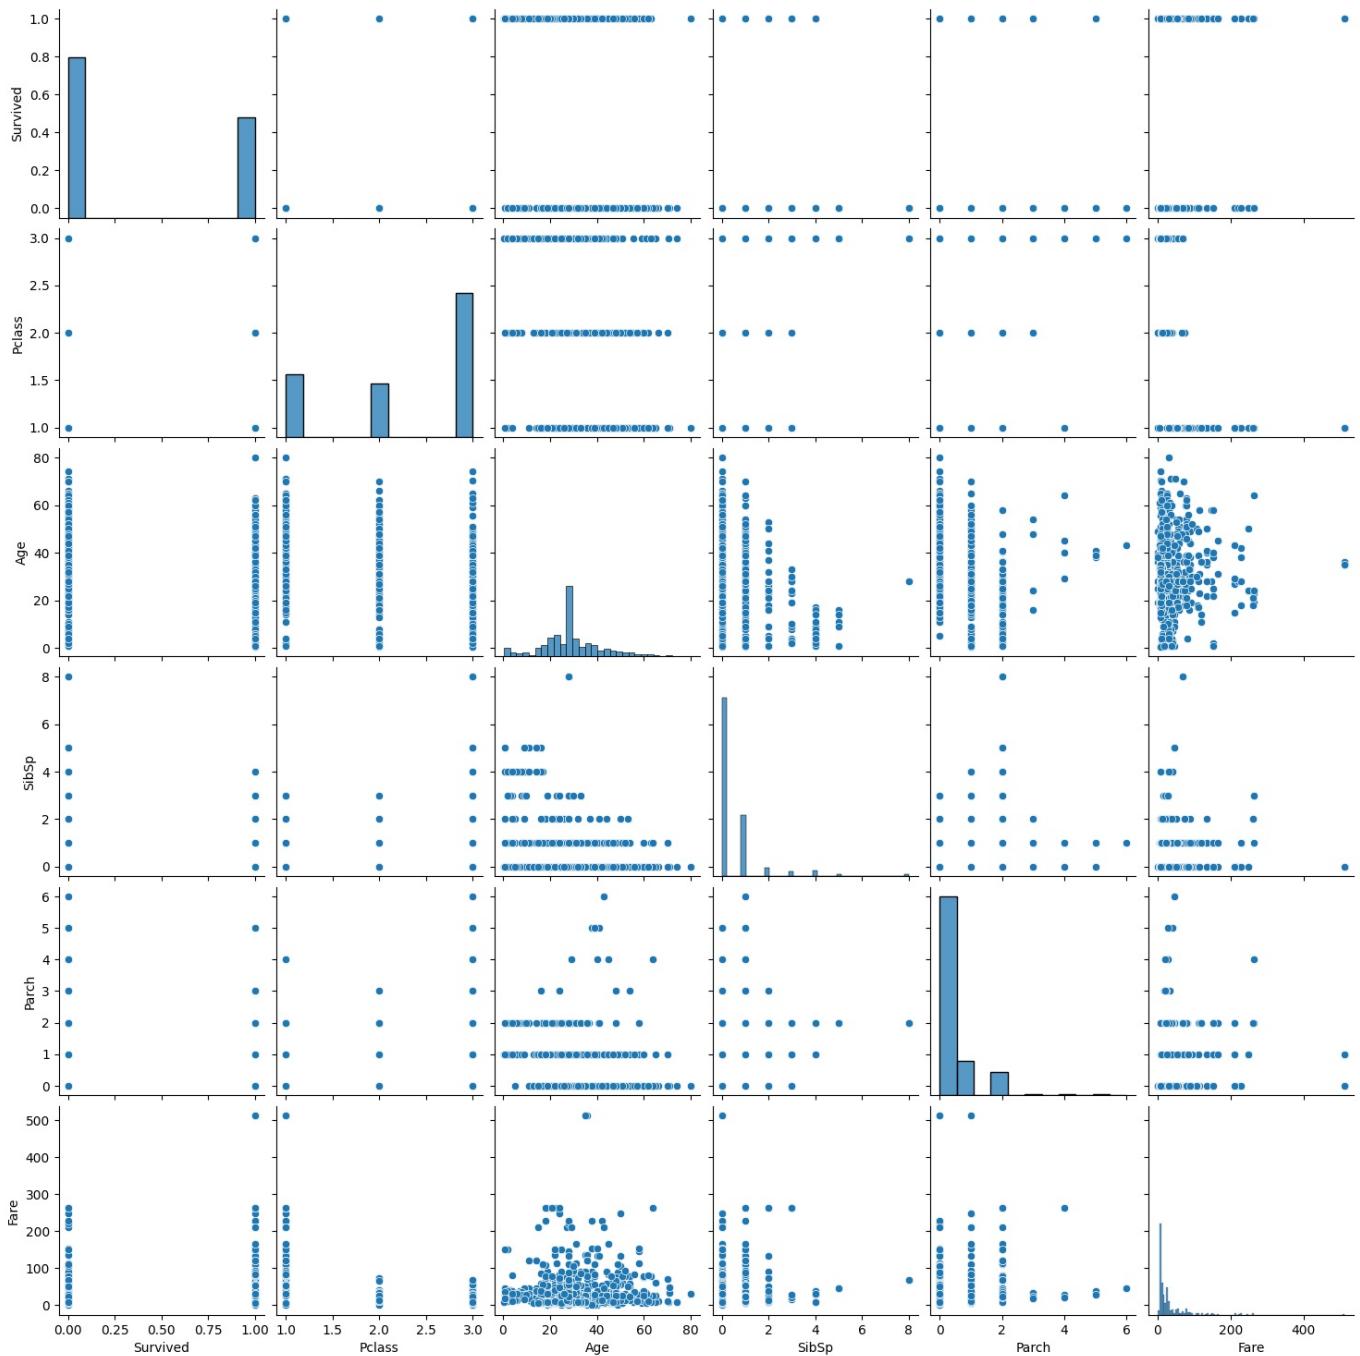


```
In [26]: fig,ax=plt.subplots(1,3,figsize=(12,5))
sns.scatterplot(x="Age",y="Fare",data=df1,ax=ax[0])
sns.scatterplot(x="Age",y="Fare",data=df1,hue="Sex",ax=ax[1])
sns.scatterplot(x="Age",y="Fare",data=df1,hue="Sex",style="Age_group",ax=ax[2])
ax[0].set_title("Scatter plot of Age and Fare",color="Blue")
ax[1].set_title("Age vs Fare Based on Gender ",color="Blue")
ax[2].set_title("Age vs Fare Based on Gender and Age Group ",color="Blue")
plt.show()
```



```
In [27]: sns.pairplot(df1)
```

```
Out[27]: <seaborn.axisgrid.PairGrid at 0x75a3fa724680>
```



```
In [28]: num_col=[ "Survived", "Pclass", "Age", "SibSp", "Parch", "Fare"]
num_col
```

```
Out[28]: ['Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare']
```

```
In [29]: cor=df1[num_col].corr()  
cor
```

	Survived	Pclass	Age	SibSp	Parch	Fare
Survived	1.000000	-0.338481	-0.064910	-0.035322	0.081629	0.257307
Pclass	-0.338481	1.000000	-0.339898	0.083081	0.018443	-0.549500
Age	-0.064910	-0.339898	1.000000	-0.233296	-0.172482	0.096688
SibSp	-0.035322	0.083081	-0.233296	1.000000	0.414838	0.159651
Parch	0.081629	0.018443	-0.172482	0.414838	1.000000	0.216225
Fare	0.257307	-0.549500	0.096688	0.159651	0.216225	1.000000

```
In [30]: sns.heatmap(cor,annot=True)
```

```
Out[30]: <Axes: >
```

